

Introduction to High performance computing using the CHPC

Werner Smidt

Outline

- Introduction to modern computing
- Command-line interface explained
- Running BLAST
- Preparing scripts for the CHPC
- Other usage examples

Modern computing

- Computing power has expanded over the last few decades
- Possible to run large bioinformatic applications
- Availability of applications
 - Proprietary
 - Open source
- Local and remote applications
 - BLAST example of both
- GUI interfaces
 - CLC
 - NCBI tool/EBI etc

Modern Computing

- Scalability
- Common for PCs to have multiple CPUs
- Faster completion of tasks
- Larger storage
- More RAM for heavy-duty tasks
- Local PCs -> Servers -> HPC



Limitations of local PCs

- Modern personal computers are powerful, but ...
 - Still struggle with certain tasks
- For instance:
 - Genome assemblers - Velvet for a ~300MB genome uses ~48GB RAM
 - Read mapping possible, but takes a long time
 - Large phylogenies

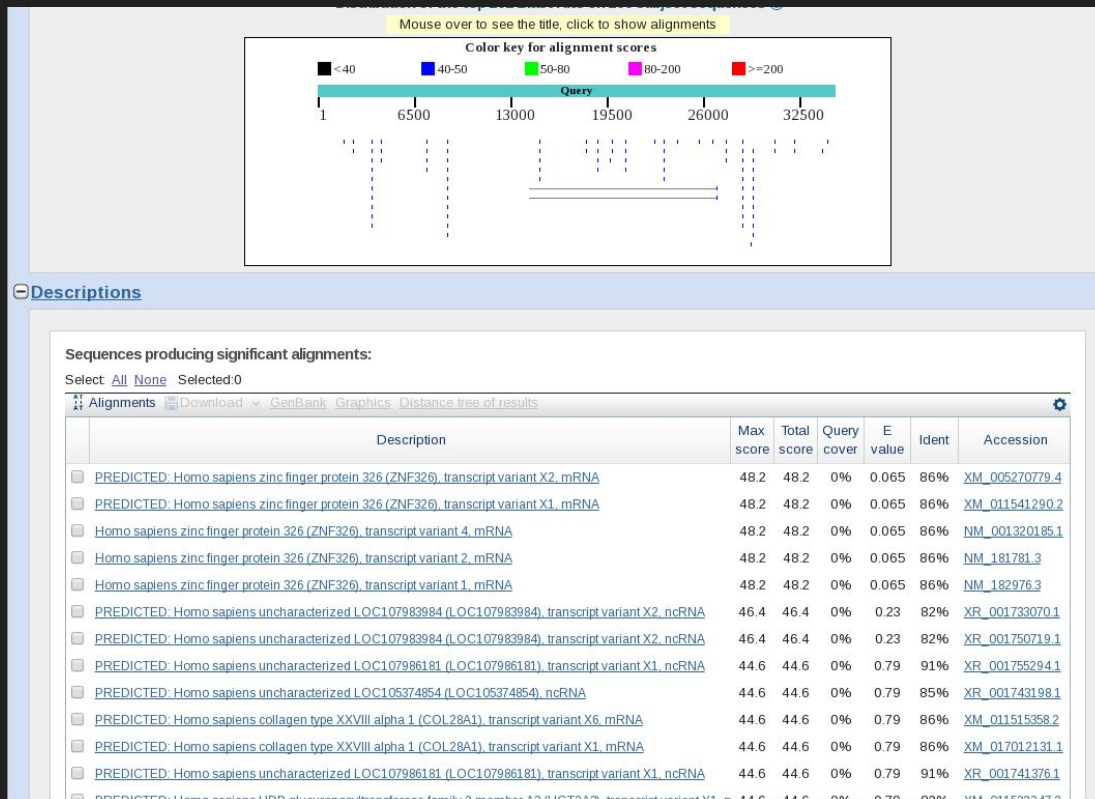
How did I get into this (nut)shell?

- Any sort of interface to the underlying operating system is a “shell”
 - Graphical
 - Text based
- Here, “shell” will exclusively pertain to the textual component
- Many BIF programs have textual interfaces:
 - Genome assemblers
 - Read mappers
 - Phylogenetic software
 - etc.

But ... why?

- Hard to make a graphical interface suitable for all the tools you may need
- A shell, in the textual form, enforces a standard
- A form of “cognitive ergonomics” or ... whatever
- Allows easy creation of ‘scripts’ (think experimental protocol):
 - Sequential command execution
 - Align -> Trim -> Motif discovery
 - Quality filtering/trimming -> Assembly phase 1 -> Assembly phase 2 -> Annotation
- Personally I like to stay away from the “P” word ...
- Easier/more reliable to connect to a remote textual shell

An example



Command line BLAST

Command line BLAST (via SSH)

- Secure SHell is a tool to connect to a remote server
- This happens within the same terminal
- But logging in to another server means:
 - All the commands are executed on that remote server
 - You have access to software on that server
 - When accessing storage, you are accessing the *remote* storage (from within SSH)
 - But you can upload your data to the relevant server

Command line BLAST via SSH

Threads example

```
[~/Sources/git/miscguides/chpc/blast]
werner.local-> time blastn -num_threads 1 -query test2.fsa -db blast/rna -word_size
9 > /dev/zero

real    0m4.101s
user    0m4.092s
sys     0m0.008s
[~/Sources/git/miscguides/chpc/blast]
werner.local-> time blastn -num_threads 2 -query test2.fsa -db blast/rna -word_size
9 > /dev/zero

real    0m2.091s
user    0m4.096s
sys     0m0.004s
[~/Sources/git/miscguides/chpc/blast]
werner.local-> time blastn -num_threads 3 -query test2.fsa -db blast/rna -word_size
9 > /dev/zero

real    0m1.437s
user    0m4.040s
sys     0m0.004s
[~/Sources/git/miscguides/chpc/blast]
werner.local-> time blastn -num_threads 4 -query test2.fsa -db blast/rna -word_size
9 > /dev/zero

real    0m1.194s
user    0m4.046s
sys     0m0.013s
[~/Sources/git/miscguides/chpc/blast]
werner.local->
```

The Centre for High Performance Computing

- It's massive
- There are 1368 compute nodes
 - Equates to 32,832 cores (workers)
- Massive storage capacity
 - 4 PB = 4,000 TB = 4,000,000 GB
- Each node typically has 128GB RAM and 24 cores



The Centre for High Performance Computing

- You can imagine that managing such a resource is a mammoth task
- Also, while there are thousands of nodes . . .
 - Ridiculous to log into each node to see which one is free
 - Better to have a queuing manager that can handle jobs you would like to run
- Because of this:
 - Commands should not be run directly
 - Instead, they need to be “wrapped” in a script that can:
 - Tell the queuing manager what is needed i.t.o resources
 - Who is running the job (group)
- It also safeguards against “hogging” the resource
 - Lower priority to too-frequent users ensuring a fair usage policy

The Centre for High Performance Computing

- Most BIF application you would like to run are not available implicitly
- The CHPC uses the GNU Modules suite to handle this
- Allows simultaneous existence of incompatible software/different versions
- Makes the task of adding new applications easier

[~]

werner.local-> ssh wsmidt@lengau.chpc.ac.za

Last login: Thu Jun 15 06:45:55 2017 from zoidberg.bi.up.ac.za

Welcome to LENGAU

[~]

wsmidt-> tophat --version

-bash: tophat: command not found

[~]

wsmidt-> module add chpc/BIOMODULES

[~]

wsmidt-> module load tophat/2.1.1

[~]

wsmidt-> tophat --version

TopHat v2.1.1

[~]

wsmidt-> module avail 2>&1 | head -n 4

----- /apps/chpc/scripts/modules/bio/app -----

ABBySS/1.9.0

ABBySS/2.0.2

[~]

wsmidt-> █

The BLAST example (again)

- So now we'll do something a little more fancy
- We'll run BLASTX against the refseq_protein database
- Knowing it would be big task
 - Does translation for the entire sequence in different reading frames and on both strands
 - Tries to find matching proteins
- Perfect opportunity to try out the CHPC
- As stated before, we need to wrap this command in a script (PBS script)
- Before we get too excited, let's examine the final PBS script we'll use

BLASTX PBS script

- Let's look at a typical (and trivial) PBS script

```
#!/bin/bash
#PBS -l select=1:ncpus=24:nodetype=haswell_reg
#PBS -P CBB10905
#PBS -q smp
#PBS -l walltime=48:00:00
#PBS -o /home/wsmidt/outputs/myblast.out
#PBS -e /home/wsmidt/outputs/myblast.err

#PBS -N mytestblast
#PBS -M werner.smidt@gmail.com

cd $PBS_O_WORKDIR
export BLASTDB="/mnt/lustre/bsp/NCBI/BLAST"

module add chpc/BIOMODULES
module load ncbi-blast-2.6.0

blastx -query test.fsa -db refseq_protein -num_threads 24 #The actual command
```

BLASTX PBS script

- Yikes!
- OK, so this is the command we would like to run

```
blastx -query test.fsa -db refseq_protein -num_threads 24
```

- To do this, we need to set up a few things first
 - Load the appropriate modules
 - Do additional setup procedures in order for the BLAST to run
 - Specify where screen output should go
 - Tell the queue manager the resources we require

Let's work a little in reverse

- Let's look at the lines preceding the blastx command we would like to run

```
module load chpc/BIOMODULES  
module load ncbi-blast-2.6.0
```

- The command “module” is used to set up our shell to execute the applications we want
- Modules are divided into categories
 - BIF tools are under the chpc/BIOMODULES “category”
 - Now the BLAST suite of utilities can be loaded
-

What do I need?

- We have everything in place to run the command
- Still need to set up the script so it can be added (submitted) to the queue

```
#!/bin/bash
```

```
#PBS -l select=1:ncpus=24:nodetype=haswell_reg
```

```
#PBS -P CBBI0905
```

- `#!/bin/bash` : We would like to use the “BASH” shell
- `PBS -l`: sets multiple options
 - Select 1 node
 - Gimme ncpus=24 cpus on the new shiny haswell_reg nodes
- `PBS -P`: I'll be submitting this under project CBBI0905

...continued

```
#PBS -q smp
```

```
#PBS -l walltime=48:00:00
```

- We would like to use the 'smp' queue.
- We expect the job to run no longer than 48 hours
 - You cannot expect things to run forever
 - There is a 96 hour limit on this queue
 - You can make a guess or max it out if you are unsure

Where will the output go?

```
#PBS -o /home/wsmidt/outputs/myblast.out  
#PBS -e /home/wsmidt/outputs/myblast.err
```

- BLAST has output options, but I didn't specify any
- By default, the output will be written to the screen
- Since the job we're submitting will (eventually) run on a random node
 - How do we get the textual output?
- Two 'streams of text':
 - stdout - Normal output of a program
 - stderr - Additional info/error messages
- Can separate the two

Queues

Queue Name	Max. cores	Min. cores	Max. jobs		Max. time	Notes	Access
	per job		in queue	running	hrs		
serial	24	1	???	???	48	For single-node non-parallel jobs.	
smp	24	1	20	10	96	For single-node parallel jobs.	
normal	240	48	20	10	48	The standard queue for parallel jobs	
large	2400	264	10	5	48	For large parallel runs	<i>Restricted</i>
bigmem	280	28	4	1	48	For the large memory (1TiB RAM) nodes.	<i>Restricted</i>
vis	24	1	1	1	3	Visualisation node	
test	24	1	1	1	3	Normal nodes, for testing only	

Naming and notification

```
#PBS -N mytestblast  
#PBS -M werner.smidt@gmail.com  
#OPTIONAL: PBS -m abe
```

- We give the job a name that we can identify
- Ask any notification to be sent to my email address
 - By default, it will only send you an email when the job fails
- I didn't include the "optional" line:
 - Mails can be sent at the (b)eginning, (e)nd and when a job is (a)borted

BLASTX PBS script

- Let's look at a typical (and trivial) PBS script

```
#!/bin/bash
#PBS -l select=1:ncpus=24:nodetype=haswell_reg
#PBS -P CBB10905
#PBS -q smp
#PBS -l walltime=48:00:00
#PBS -o /home/wsmidt/outputs/myblast.out
#PBS -e /home/wsmidt/outputs/myblast.err

#PBS -N mytestblast
#PBS -M werner.smidt@gmail.com

cd $PBS_O_WORKDIR
export BLASTDB="/mnt/lustre/bsp/NCBI/BLAST"

module add chpc/BIOMODULES
module load ncbi-blast-2.6.0

blastx -query test.fsa -db refseq_protein -num_threads 24 #The actual command
```

Bear with me . . .

- Assume I “somehow” got my script to the CHPC
- Assume I did everything else by the book
- To submit my job to the queue, I need to run the command

```
qsub myblast.pbs
```

Qsub

Output

What's next?

- This is just a trivial example of CHPC usage
- Things you can imagine doing:
 - Mapping
 - Genome assembly
 - Population genetics
 - Phylogenetics
 - Molecular dynamics
 - Etc
- Practical sessions?
 - Accounts?
 - I'll restrict my role to technical parts
 - Members of your group can present what they have done?